



**CS310 – AI Foundations**

**Andrew Abel**

**January 2024**

# **Week 4: Adversarial Search**

# Welcome!

---

- Welcome to CS310 week 4!
- Adversarial search
  - Introducing dealing with another agent
  - Handling small games
  - Minimax Algorithm
- Pruning Game Trees
  - Alpha-Beta pruning
  - Making your algorithm more efficient

# Adversarial Search

---

## Learning Outcomes

---

- At the end of this topic, you will know:
  - How two-player zero-sum games are an extension of the single agent problems we've examined so far
  - What a game tree looks like
  - How we can solve a game tree by assuming that the opponent plays rationally
  - How the Minimax algorithm works

# Games & Multi-agent environments

---

- Which Problems can we Solve (see week 2)?
- Task environments suitable for our search algorithms so far:
  - Fully observable (the easy option)
  - Single agent (the easy option)
  - Deterministic (the easy option)
  - Sequential (the normal option)
  - Static (the easy option)
  - Discrete (the easy option)

# Games

---

- To play a game, we need to relax the assumption that only one agent can change the state of the world.
  - Need to move away from static and deterministic search trees/graphs!
- Game theory: any environment with multiple agents in it can be regarded as a game.
- In AI, simple games are often what game theorists would call deterministic, turn-taking, two-player, zero-sum games of perfect information.
  - Examples: chess, checkers, Connect 4, Shogi, Othello, go, tic-tac-toe, dots and boxes...
  - Deterministic: no element of chance
  - Turn-taking: players take turns
  - Two player – 2 agents
  - Zero sum – 1 person wins, other loses. Can't both win
  - Perfect information – Everybody has all information about state of the game

## Features of these Games

---

- Fully observable: game state is visible to both players
  - Perfect information
- Deterministic: no element of chance
- Sequential: action taken now affects future choices
- Static: the world doesn't change during deliberation
- Discrete: in each game state, there are a finite number of moves
- Known: the outcomes of an action are fully defined
- **Multi agent**: in the search, we have to allow for the fact that our opponent can also affect the game state when it is their turn, and will be planning against us

# Min-Max Game

---



# Representing a Game

---

- We are going to deal with a game where we have 1 opponent, so a 2 player game that is perfect information, turn based and zero sum
  - This is a min-max game
- We need:
  - An initial state, including who plays first
  - A successor (state) function, like our Nextstates() python function
  - A terminal state/test which determines whether a given state is a end state of the game (i.e. the game is over)
  - A **utility function** - for terminal states only, this is the reward each player gets (e.g. +1 for win, -1 for lose)
  - In a zero-sum game, the wins and losses for the two players sum up to zero (e.g. one player wins £7, the other one loses £7)

## The Two Players: Max and Min

---

- Let's call the two players Max and Min: Max goes first
- Max's task is to maximise Max's reward (i.e. win!)
- Min's task is to minimise Max's reward (and therefore maximise Min's reward)
- Let's say Max can take actions  $a$ ,  $b$  or  $c$  - which one will give Max the best reward when the game ends?
- To answer that question, we need to explore the game tree to a sufficient depth, and assume that Min plays optimally to minimise the reward that Max gets
  - The game tree is like a search tree, but different layers belong to different agents
  - Cannot just look for one path, because we need cooperation of an opponent

# An Example

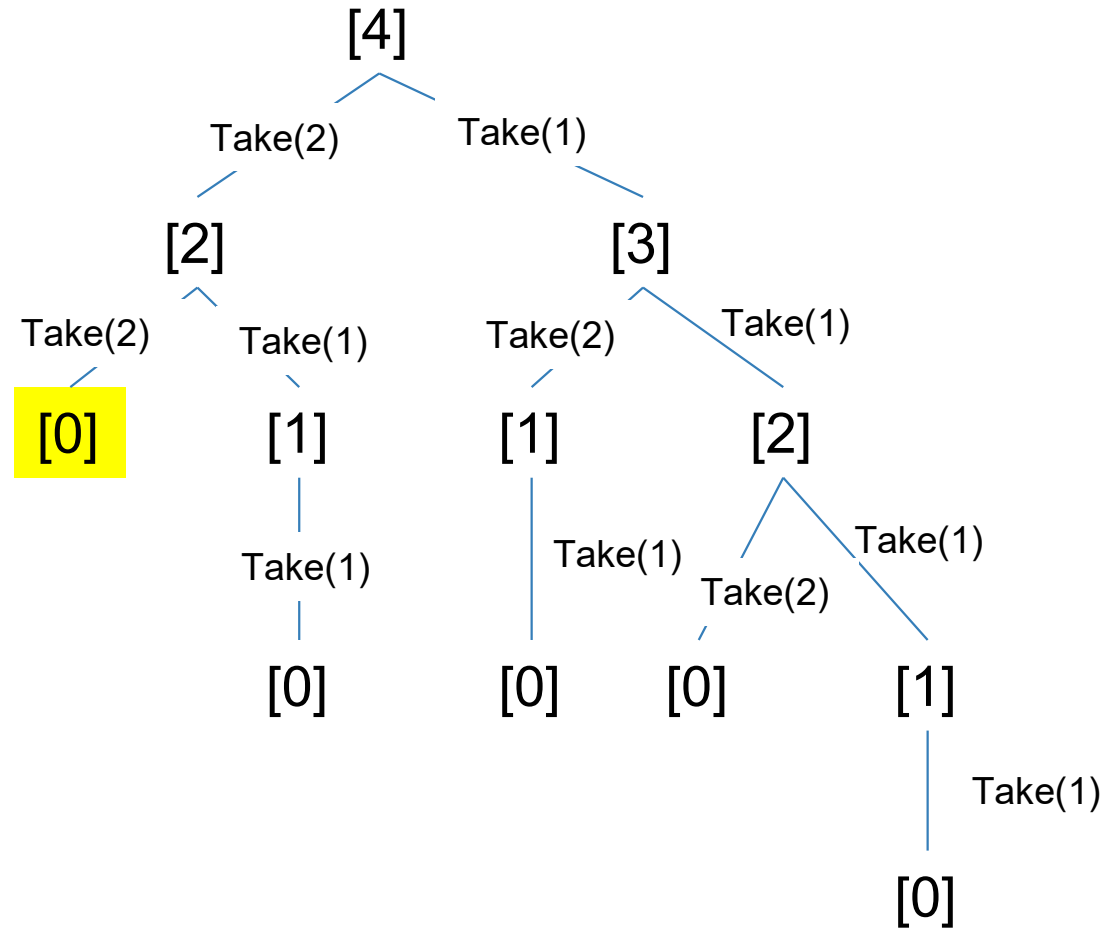
---

- The world's least fun game:
  - Four coins in a row, each player can pick up one coin or two coins.
  - The player who picks up the last coin wins.
  - Max plays first. What move should Max make?
- Initial State: Table with 4 coins
- Actions: take(1), take (2)



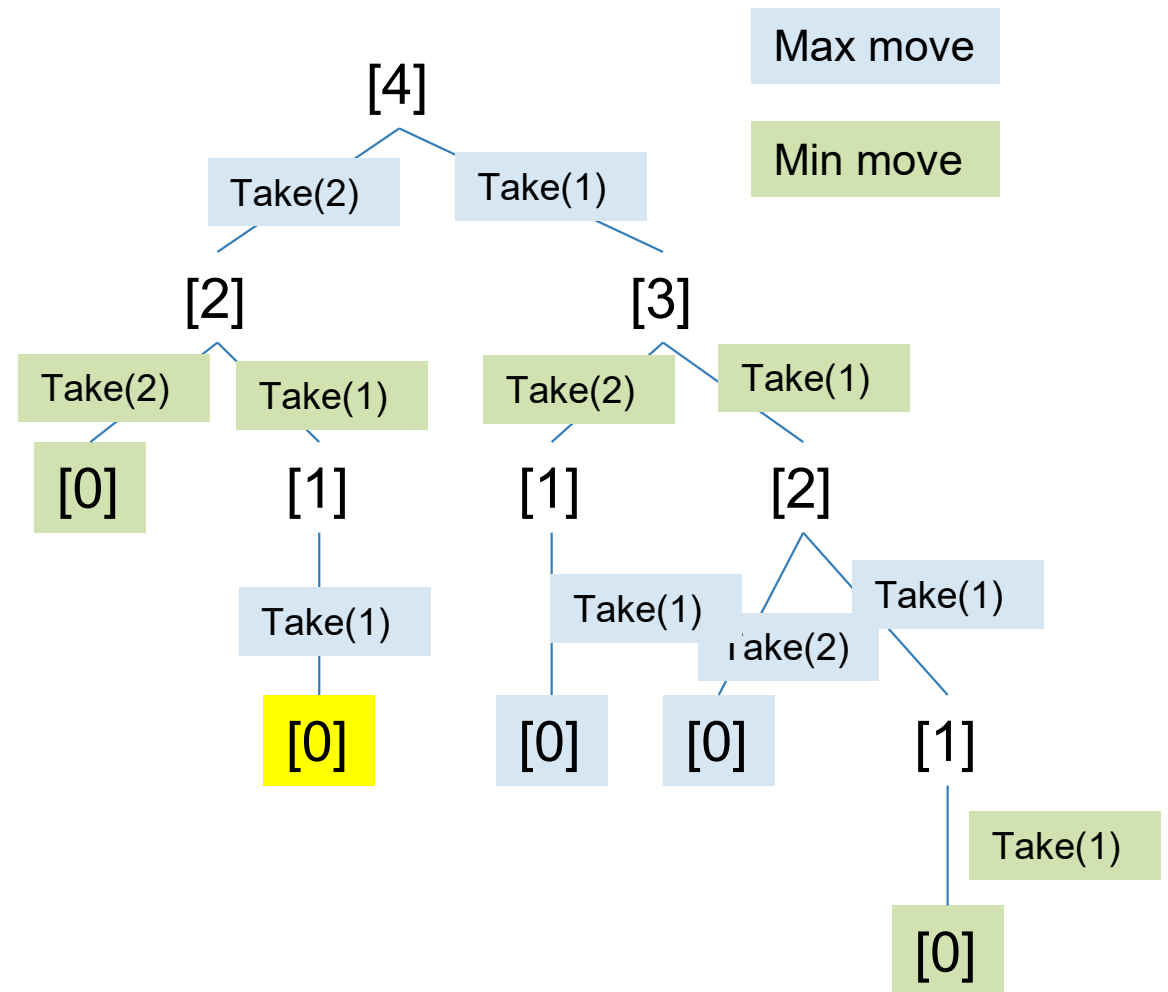
## Considering as a Search Tree

- Building a search tree of possible moves
- Immediate optimal solution on left of tree, would find first with breadth first or depth first!
- May even find first with heuristics
- But, this is a game tree, not a search tree, so we need to consider differently



# Considering as a Game Tree

- Different layers belong to different agents
- Each player makes an action
- Terminal nodes, [0]
- Utility function (reward)
  - +1 depending on player
- Possible optimal search path to victory through  
max(2), min(2), max(1)
- But playing optimally, min will never do that! So its not a sensible winnable path.



# MiniMax Algorithm

---

# The Minimax Algorithm

---

- Should allow us to find an optimal solution for a 2 player zero sum game with perfect information
- A structured way of solving the problem introduced in previous slides
- Considers the tree as a game tree rather than a search tree
- Computes the payoff on the assumption that both players play the optimal strategy

# The Minimax Algorithm

---

- The Minimax Value
  - The minimax value of a node is the utility of the node if the node is a terminal (i.e. the reward)
  - If the node is a non-terminal Max node, the minimax value of the node is the maximum of the minimax values of all of the node's successors
  - If the node is a non-terminal Min node, the minimax value of the node is the minimum of the minimax values of all of the node's successors
  - Recursive definition: results in a depth-first traversal of the game tree
    - Start at the terminal, work your way backwards
    - Also known as Backwards Induction




# The Minimax Algorithm

---

- `MaxValue(state)` returns a utility value
  - if `Terminal-Test(state)` then return `Utility(state)`
  - $v \leftarrow \text{MinimalGameValue } (= -\infty)$
  - for `s` in `Successors(state)` do
    - $v \leftarrow \text{Max}(v, \text{MinValue}(s))$
  - return `v`
- `MinValue(state)` returns a utility value
  - if `Terminal-Test(state)` then return `Utility(state)`
  - $v \leftarrow \text{MaximalGameValue } (= +\infty)$
  - for `s` in `Successors(state)` do
    - $v \leftarrow \text{Min}(v, \text{MaxValue}(s))$
  - return `v`
- If it's a terminal node, return the reward (utility)
- If it's not, initially, not knowing the details, we can assume the worst for the player
- Can either calculate min value, or assume  $-\infty$
- For successors, we compute the min values, and then take the max of the min value
- Vice versa for `minValue`

# The Minimax Algorithm: Step-By-Step

---




Carry out maxVal function  
Want to know the minimum  
possible value

Still undecided as not a terminal  
known  
Have to consider successors

# The Minimax Algorithm: Step-By-Step


---



Min value of first successor  
Again, unknown, so max value,  
or + infinity  
Need to look at successors

# The Minimax Algorithm: Step-By-Step

---




Terminal node located!

Apply max value (its max's turn!)  
Terminal node, so we get a value

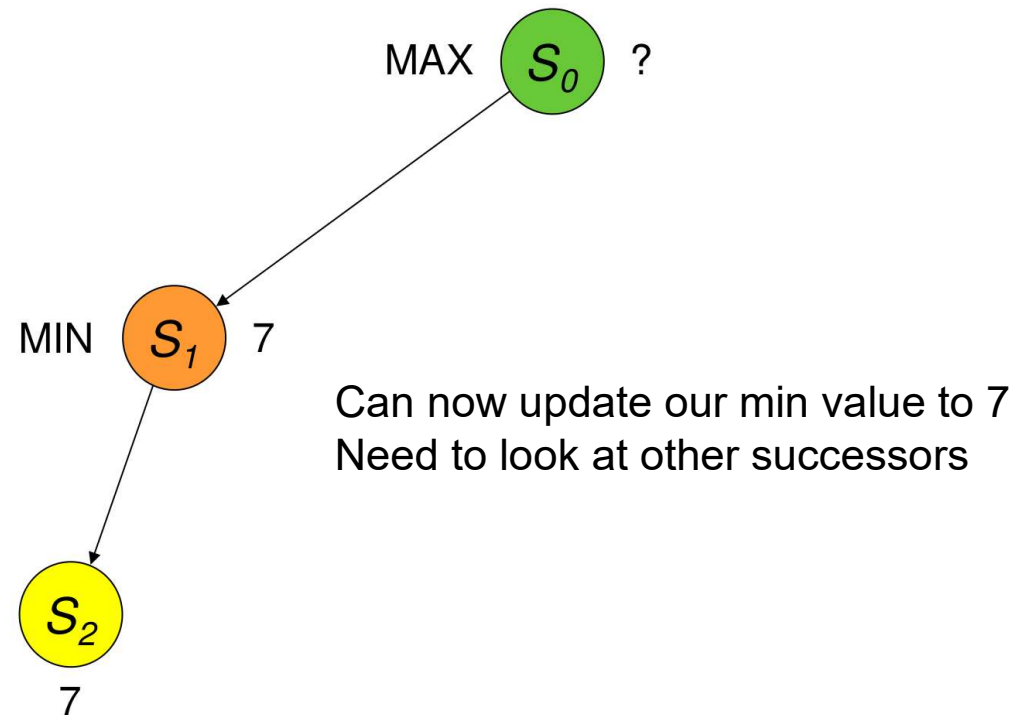
# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---



# The Minimax Algorithm: Step-By-Step


---






# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---



# The Minimax Algorithm: Step-By-Step

---




Max stays 5  
Max of mins  
 $\text{Max}(5, 2) = 5$

Min is 2 (min of max)  
 $\text{Min}(2, 8) = 2$




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---



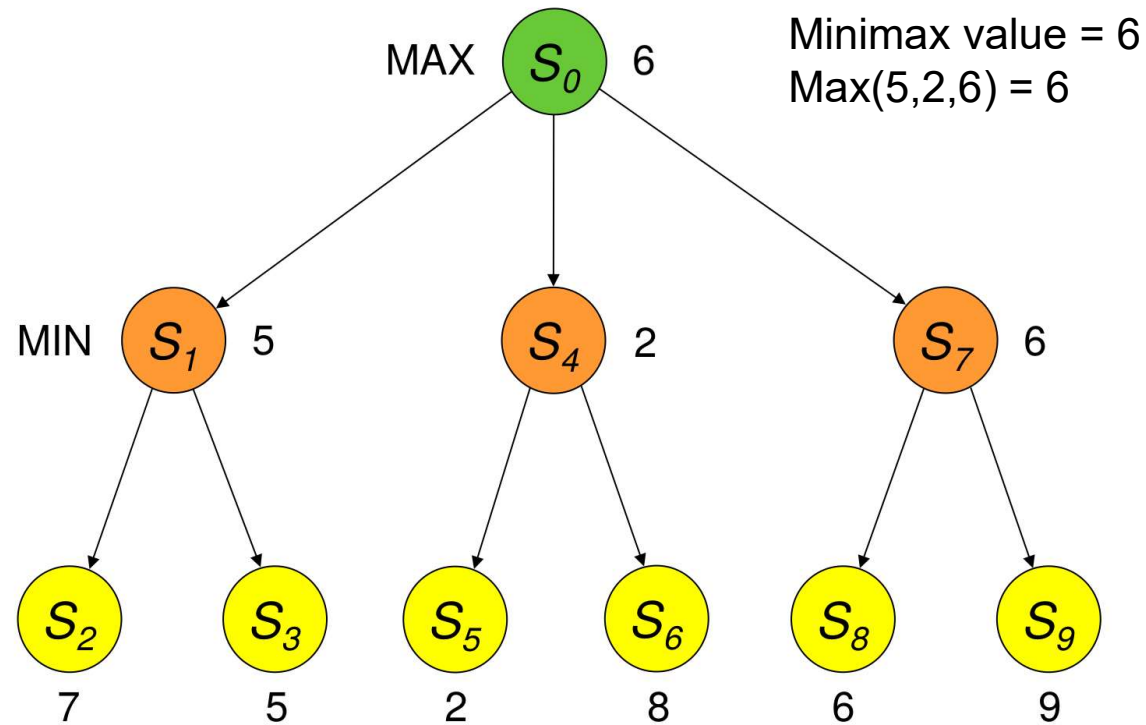
# The Minimax Algorithm: Step-By-Step

---




# The Minimax Algorithm: Step-By-Step

---



# The Minimax Algorithm: Step-By-Step

---



How to find optimum path?  
Follow all the '6' nodes  
Go to successor

This maximises max reward  
(6 or 9)  
Assuming min also plays  
optimally

Min can only go to  $S_8$  to  
minimise max's reward

Its not the highest reward, but it is the  
highest reward assuming that both players  
play optimally!

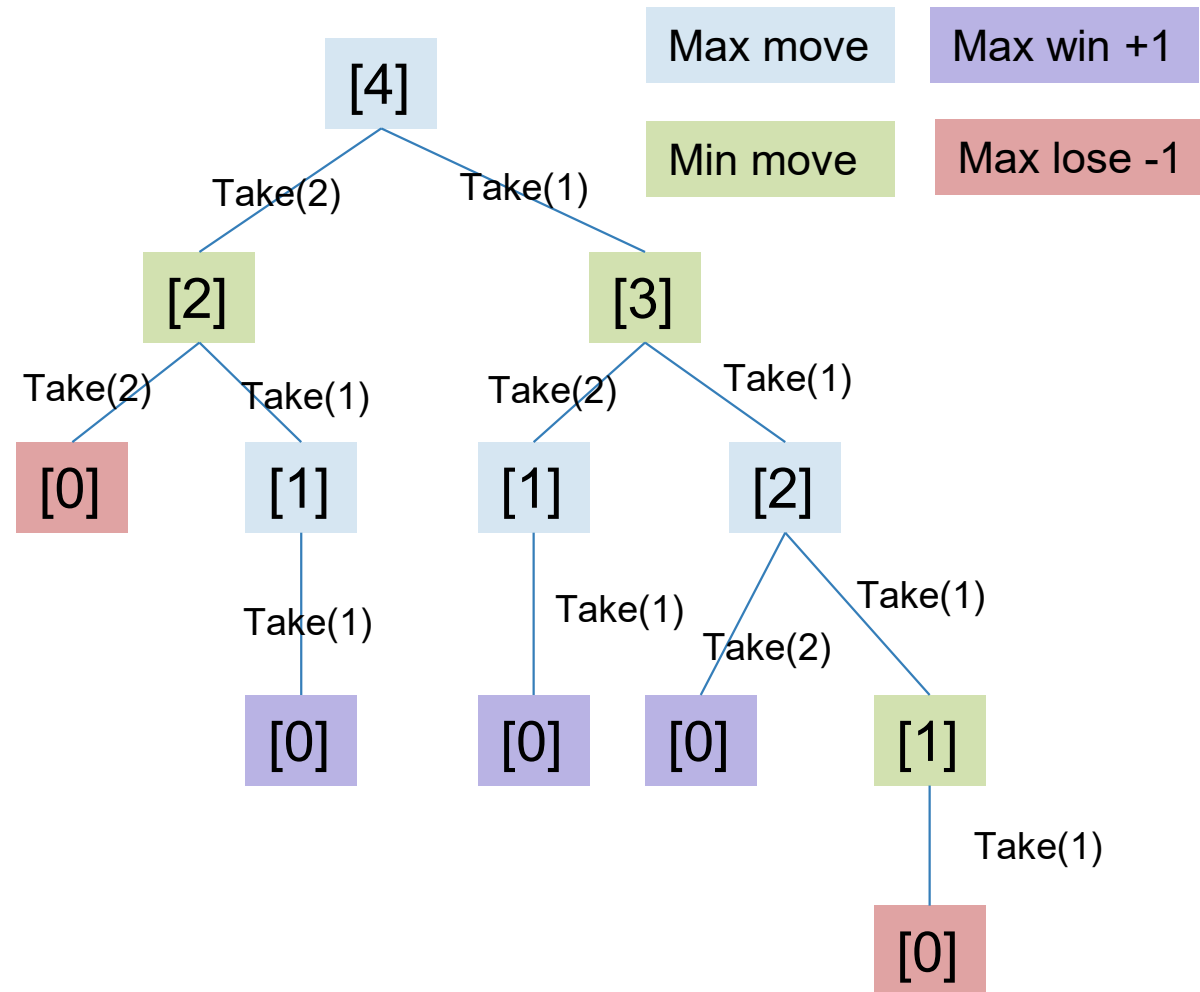


# Coin Game

---

# Our Coin game

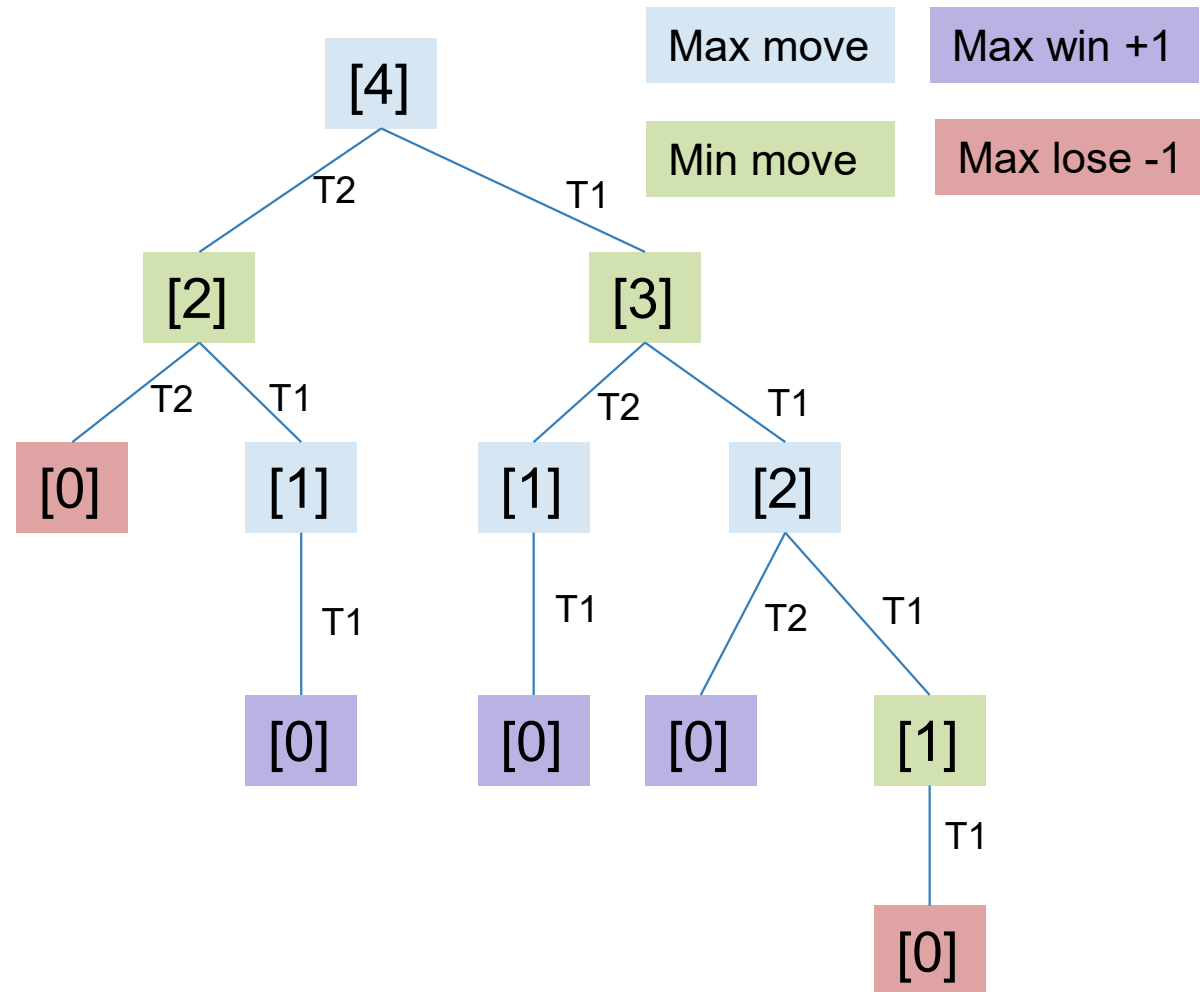
- Different layers belong to different agents
- Each player makes an action
- Terminal nodes, [0]
- Utility function (reward)
  - +1 for max win (purple)
  - -1 for min win (red)
- Start with max



# Our Coin game

---


- Left path first



## Our Coin game

---

- Left path first
- Find terminal state
- First state, so set val to -1



Max move

Max win +1


Min move

Max lose -1

# Our Coin game

---

- Other successor
- Finds +1
- Sets val to +1



Max move


Max win +1

Min move

Max lose -1

## Our Coin game

- As its mins move, we look to find the minimal value
- Min of -1 and +1



Max move

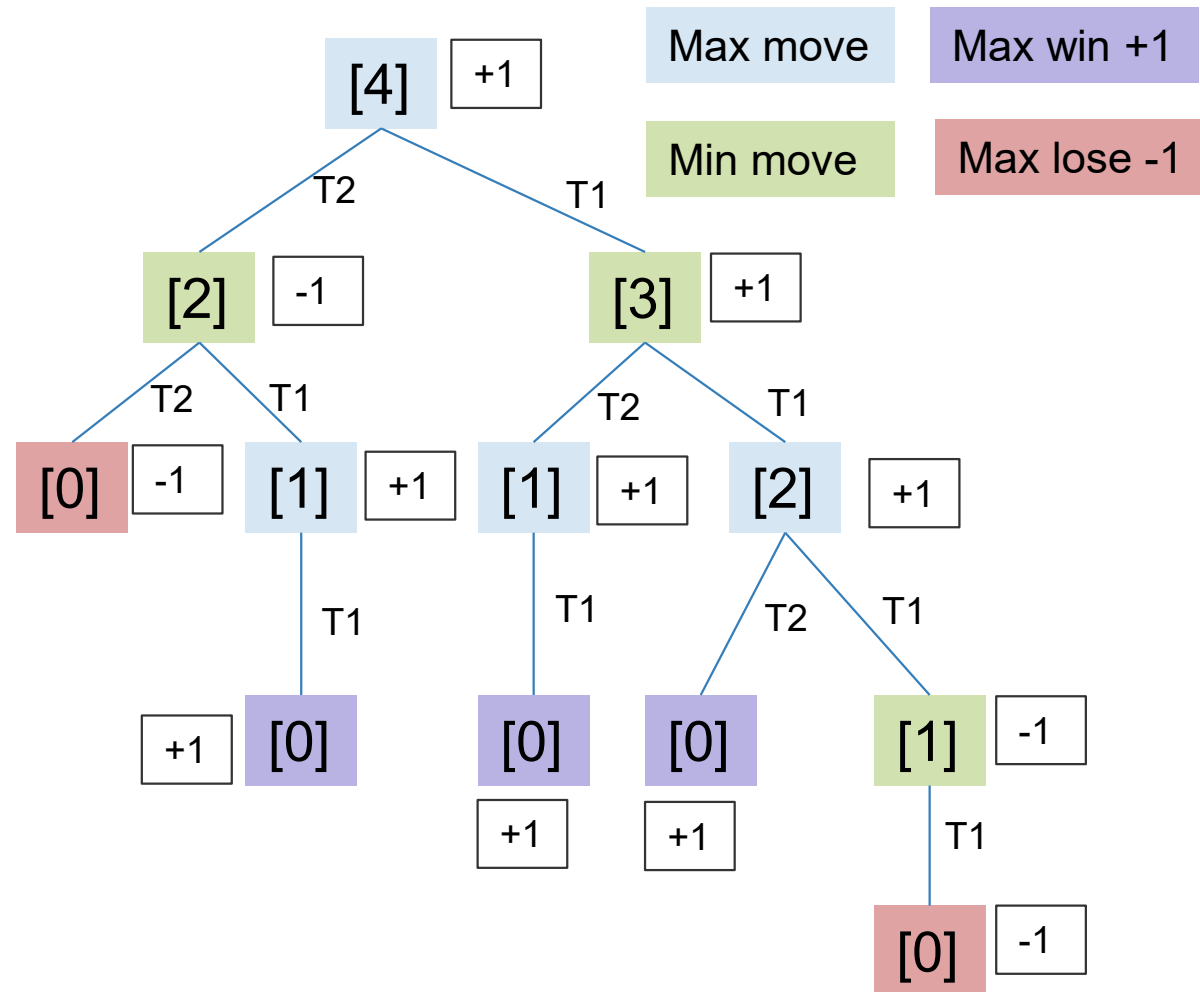
Max win +1

Min move

Max lose -1

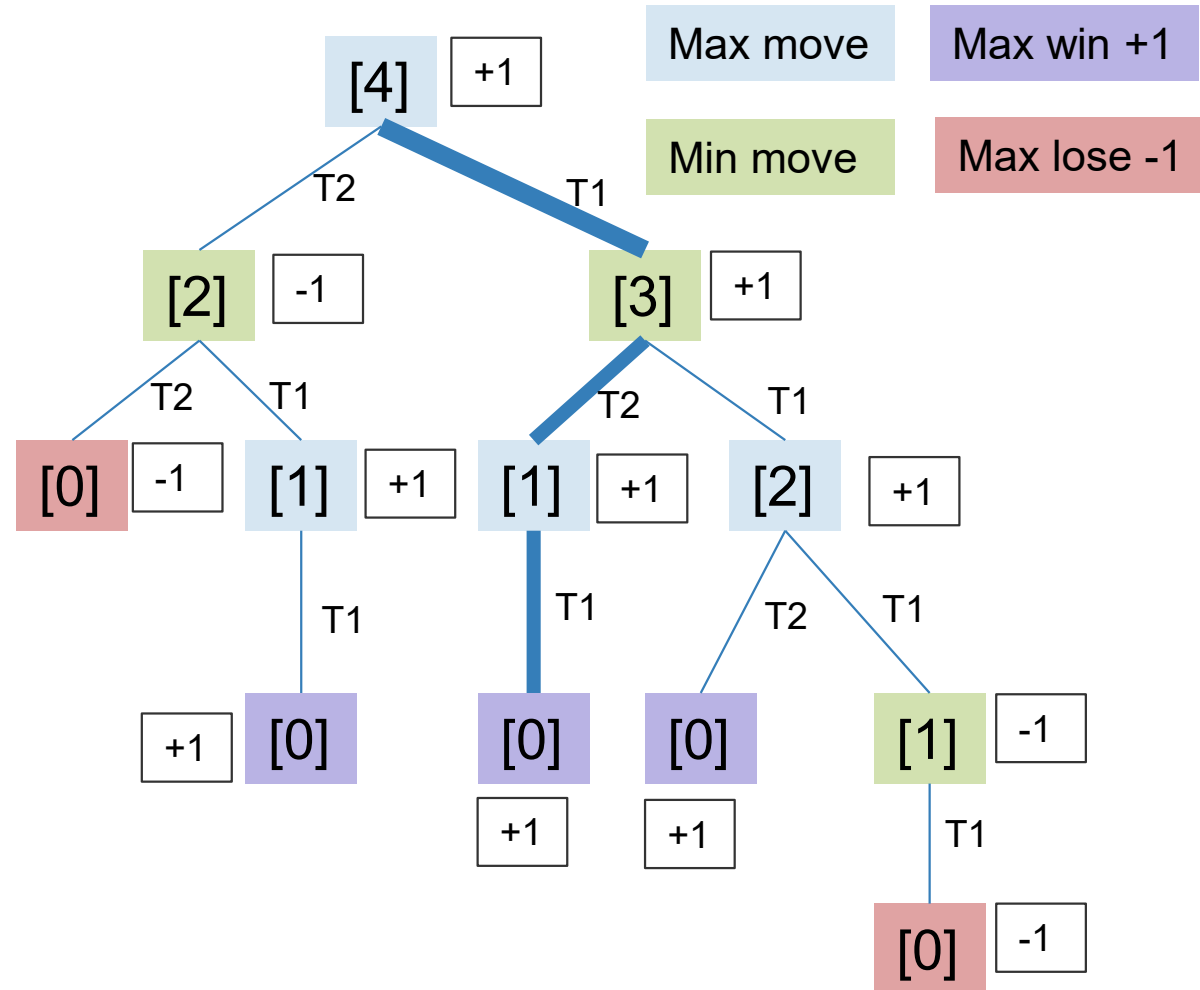
# Our Coin game

- Repeat same process
- Find terminal states and work upwards



# Our Coin game

- Advice:
- Always take 1 coin!
  - Path in blue
- May be more than one optimal path
- Can use metrics to decide (such as first one found etc.)





## Summary

---

- 2 player perfect information game
- Game trees rather than search trees
- Minimaxing
  - However, this is not efficient, as all states need to be explored
  - What if there are time constraints?
    - i.e. speed chess?