



CS310 – AI Foundations

Andrew Abel

February 2024

Week 4: Pruning

Welcome!

- Improve minimax with pruning
 - pruning a search to improve processing time
 - Refinement of minimax
 - α and β pruning


Previously

- We introduced game trees through minimax
- We call min and max alternatively
- Applies to zero sum games with turn taking
- Max is trying to maximise their reward
- Min is trying to minimise their reward
- Assumes that both players play optimally

The Minimax Algorithm

- `MaxValue(state)` returns a utility value
 - if `Terminal-Test(state)` then return `Utility(state)`
 - $v \leftarrow \text{MinimalGameValue } (= -\infty)$
 - for `s` in `Successors(state)` do
 - $v \leftarrow \text{Max}(v, \text{MinValue}(s))$
 - return `v`
- `MinValue(state)` returns a utility value
 - if `Terminal-Test(state)` then return `Utility(state)`
 - $v \leftarrow \text{MaximalGameValue } (= +\infty)$
 - for `s` in `Successors(state)` do
 - $v \leftarrow \text{Min}(v, \text{MaxValue}(s))$
 - return `v`
- If it's a terminal node, return the reward (utility)
- If it's not, initially, not knowing the details, we can assume the worst for the player
- Can either calculate min value, or assume $-\infty$
- For successors, we compute the min values, and then take the max of the min value
- Vice versa for `minValue`

The Minimax Algorithm:



How to find optimum path?
Follow all the '6' nodes
Go to successor

This maximises max reward
(6 or 9)
Assuming min also plays
optimally

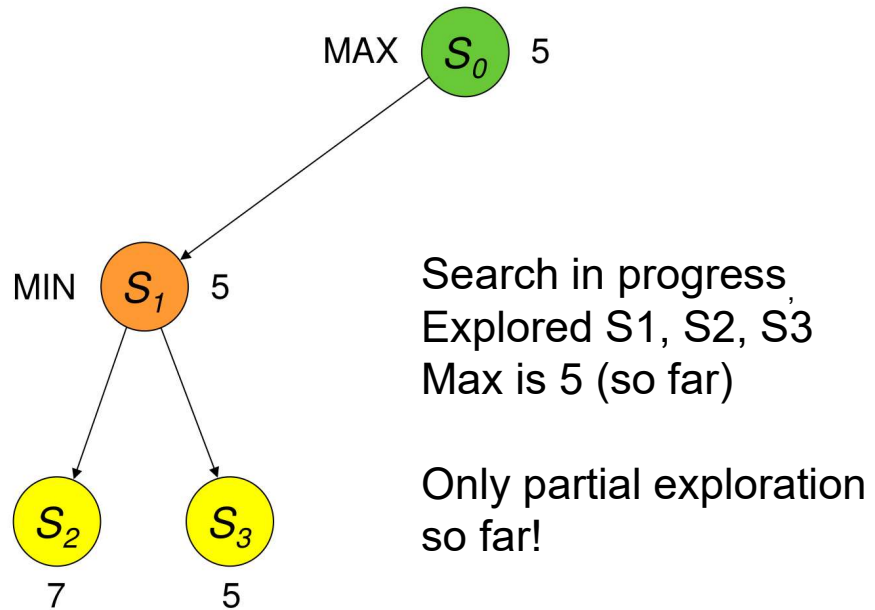
Min can only go to S_8 to
minimise max's reward

Its not the highest reward, but it is the
highest reward assuming that both players
play optimally!

Why prune the Search

- Minimax is an exhaustive search algorithm, so it is exponential in the number of moves
 - Every possible move
- Exhaustive search is not suitable
 - Even if the game takes only 20 moves, it quickly becomes unmanageable
- Even worse: inability to deal with loops/infinite plays.
 - Minimax relies on backwards induction from terminal nodes
 - How can you calculate it if there is a loop?
- We can only apply full blown minimax to very small games, or games which are close to a terminal state
- However, we can do better - sometimes we can “prune” the tree...


Pruning the Search




Min of 7 and 5 is 5

Max of 5 is 5


Pruning the Search



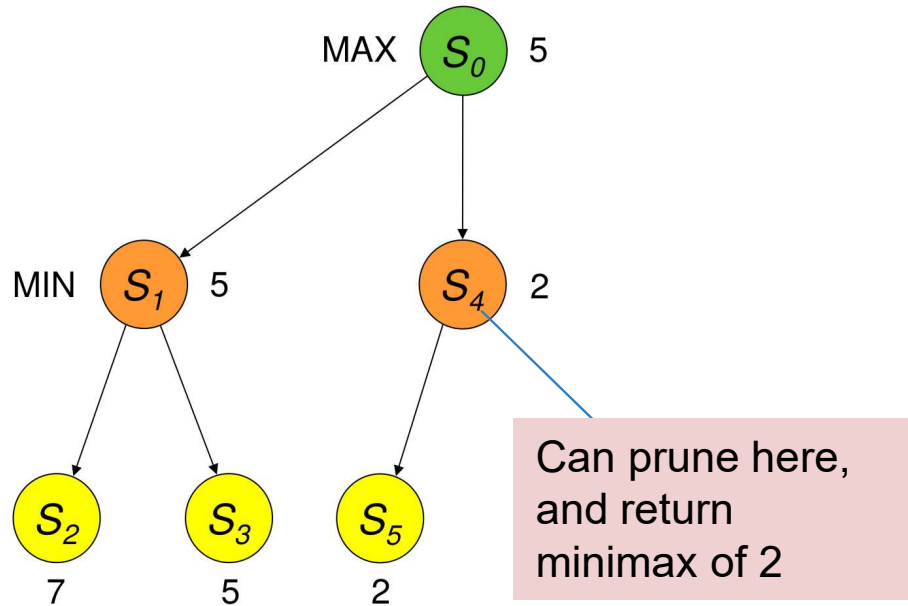
Pruning the Search



Pruning the Search



Pruning the Search



As S_5 is below a value of 5, then it means that the value of S_4 will be a minimum value of 2 (or less)

Accordingly, when S_0 picks a max value, it will pick S_1

Min now has a way to limit the damage to a max of 2.

Max has a way of going to a minimum of 5

If Max chooses S_4 , then if Min is optimal, it will always limit the reward to 2. So why would Max choose this node, regardless of what any other children are?

So Max will NEVER choose S_4

Another Pruning Example

- Eight sticks in a row. At each turn, the player can pick up three sticks, two sticks or one stick.
- The player who picks up the last stick loses.
- Example game: Max picks up 2, then Min picks up 2,
- Max picks up 1, Min picks up 2, Max picks up the last stick and thus loses.
- Max plays first. What move should Max make?
- We can prune the search once we are sure that Max has a certain win at any node

Another Pruning Example - Pointers

- If there are between 2 and 4 sticks left, you win
 - 4 left: Take 3, leave 1, other player loses
 - 3 left: Take 2, leave 1, other player loses
 - 2 left: Take 1, leave 1, other player loses
- Therefore, first move should be to have 4,3, or 2 sticks left by Max's second move
- Should move from [8] to [5]
 - i.e. pick up 3 sticks
- At [5], min has to take either 1 [4], 2 [3] or 3 [2] sticks
 - Means that max will always win

Informal Pruning

- Focus on winning positions
 - What conditions can a player have the game won?
 - Can help to remove other options from the tree
- But these are very informal!

α and β pruning

Alpha-Beta Pruning

- Alpha-Beta pruning is a method for ignoring branches of the search tree, while still finding the optimal move
 - Formalises the examples in previous slides
 - Allows us to ignore irrelevant branches of the search tree
- Given a state being explored, alpha (α) is the value of the best alternative for Max along the path to the state
 - While searching, α is the current best max number on this path
 - In other words, Max can get at least α out of the game in this part of the game tree, i.e. along current path
 - If it doesn't match α , then it can be pruned
- beta (β) is the value of the best alternative for Min along the path to the terminal state
 - In other words, Min can minimise Max's reward to be at most β in this part of the game tree
 - If min down a tree is higher than β , can be pruned

Recap: Minimax Algorithm

- `MaxValue(state)` returns a utility value
 - if `Terminal-Test(state)` then return `Utility(state)`
 - $v \leftarrow \text{MinimalGameValue} (= -\infty)$
 - for `s` in `Successors(state)` do
 - $v \leftarrow \text{Max}(v, \text{MinValue}(s))$
 - return `v`
- `MinValue(state)` returns a utility value
 - if `Terminal-Test(state)` then return `Utility(state)`
 - $v \leftarrow \text{MaximalGameValue} (= +\infty)$
 - for `s` in `Successors(state)` do
 - $v \leftarrow \text{Min}(v, \text{MaxValue}(s))$
 - return `v`

MaxValue

Takes the min value of the max values of the successors

MinValue

Takes the max value of the min values of the successors

Function Max-Value (state, α , β)

- Max-Value(state, α , β) returns a utility value
 - if Terminal-Test(state) then return Utility(state)
 - $v \leftarrow \text{MinimalGameValue}$ (initialize as $-\infty$)
 - for s in Successors(state) do
 - $v' \leftarrow \text{Min-Value}(s, \alpha , \beta)$
 - if $v' > v$, $v \leftarrow v'$
 - $v' \geq \beta$ then return v
 - if $v' > \alpha$ then $\alpha \leftarrow v'$
- return v

Function Max-Value (state, α , β)

- Max-Value(state, α , β) returns a utility value
 - if Terminal-Test(state) then return Utility(state) (same as minimax!)
 - $v \leftarrow \text{MinimalGameValue}$
 - for s in Successors(state) do
 - $v' \leftarrow \text{Min-Value}(s, \alpha, \beta)$
 - if $v' > v$, $v \leftarrow v'$ (update v if v' is better)
 - $v' \geq \beta$ then return v (**prune!** Min will prefer β)
 - if $v' > \alpha$ then $\alpha \leftarrow v'$ (update α if v' is better)
- return v

If val is greater than alpha, then set a new alpha!

If value is greater than beta, then this means a higher min value. So why would max go down this path? We can prune!

Function Min-Value (state, α , β)

- Min-Value(state, α , β) returns a utility value
 - if Terminal-Test(state) then return Utility(state)
 - $v \leftarrow \text{MaximalGameValue}$
 - for s in Successors(state) do
 - $v' \leftarrow \text{Max-Value}(s, \alpha , \beta)$
 - if $v' < v$, $v \leftarrow v'$
 - $v' \leq \alpha$ then return v
 - if $v' < \beta$ then $\beta \leftarrow v'$
- return v

If value is smaller than alpha,
then reward for max is smaller, so
max would never go down this
route

Update beta if beta is smaller

Function Min-Value (state, α , β)

- Min-Value(state, α , β) returns a utility value
 - if Terminal-Test(state) then return Utility(state)
 - $v \leftarrow \text{MaximalGameValue}$
 - for s in Successors(state) do
 - $v' \leftarrow \text{Max-Value}(s, \alpha , \beta)$
 - if $v' < v$, $v \leftarrow v'$ (update v if v' is better)
 - $v' \leq \alpha$ then return v (**prune!** Max will prefer α)
 - if $v' < \beta$ then $\beta \leftarrow v'$ (update β if v' is better)
- return v

Initial values for α and β


- we start the algorithm by calling it with $\alpha = -\infty$ and $\beta = +\infty$
- We'll have a look at detailed examples next!

Example of β -Pruning

MIN S_0 [?, ?, ?]


Min = unknown
Alpha = $-\infty$
Beta = $+\infty$

Example of β -Pruning




Started with min, so call max
Values unchanged

Example of β -Pruning




Solution found!
Can update some values

Example of β -Pruning

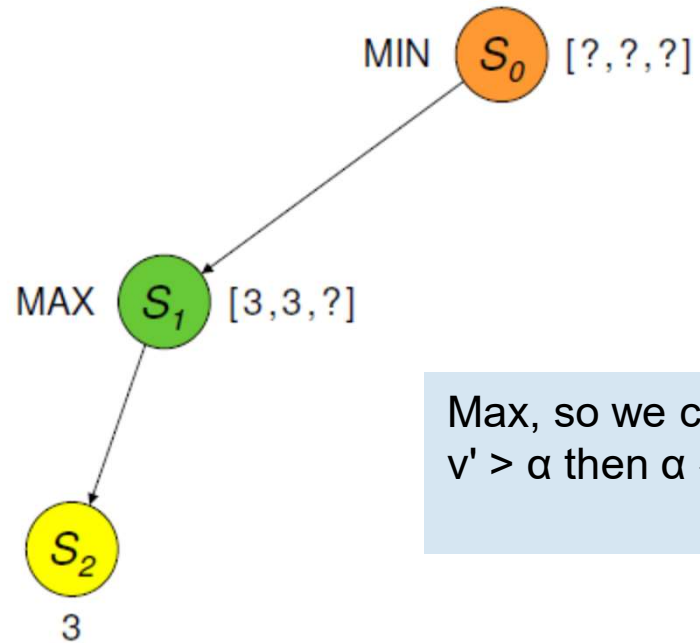


Example of β -Pruning




We update current value

Example of β -Pruning




Max, so we can update alpha
 $v' > \alpha$ then $\alpha \leftarrow v'$


Example of β -Pruning



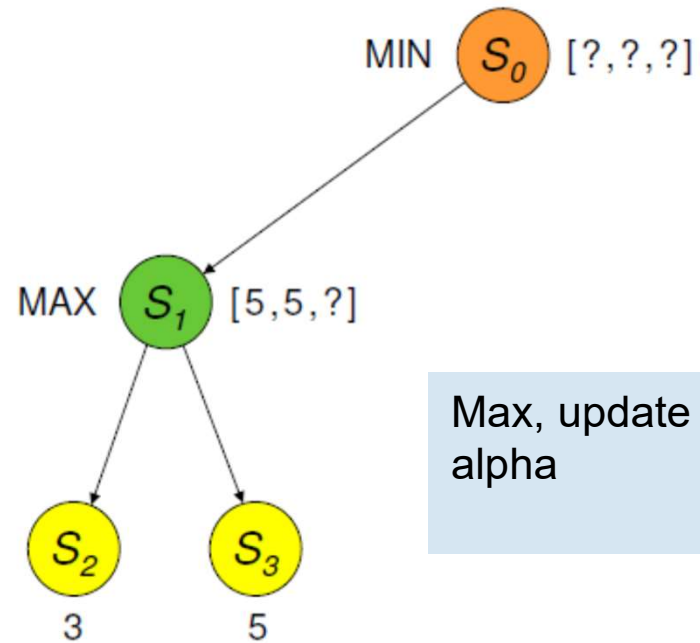
Example of β -Pruning



Example of β -Pruning




Example of β -Pruning




Max, update max value, and also alpha

Example of β -Pruning




Min value of this branch is 5

Example of β -Pruning




Can update beta

Example of β -Pruning




Improved beta can travel down path

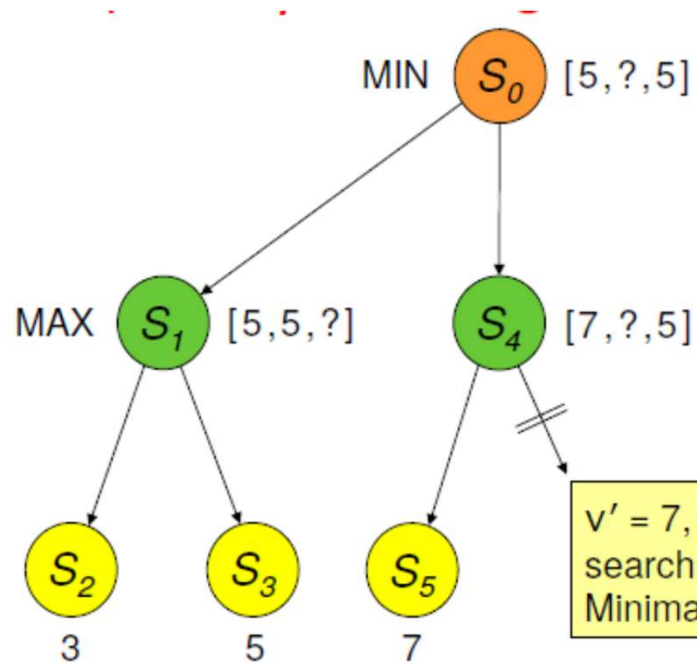
Example of β -Pruning



Example of β -Pruning



Example of β -Pruning



7 is bigger than beta, so no need to continue

Choices are 7 and 5, min of maxes is still 5

$v' = 7$, $\beta = 5$, so prune the search and return 7 as the Minimax value of S_4


Example of α -pruning

- Similar, except start with max
 - i.e. max goes first

MAX  [?,?,?]

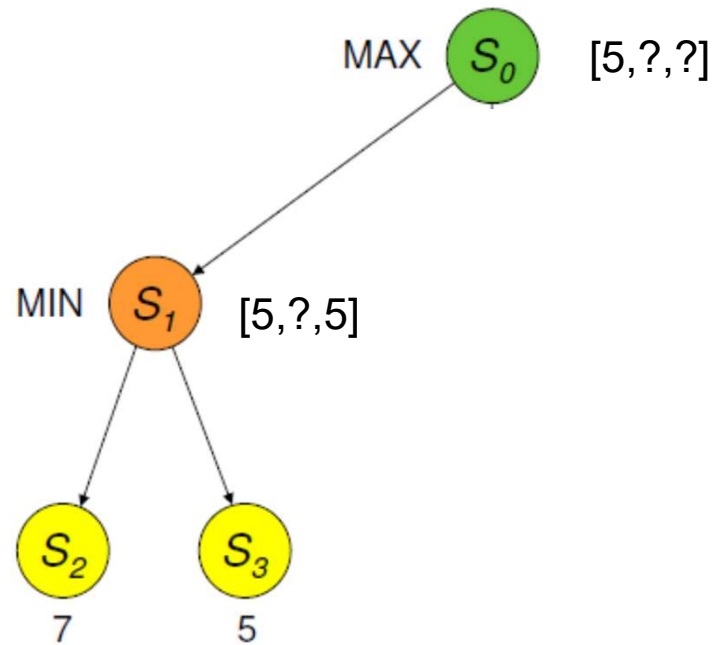
Example of α -pruning

- First path, update beta




Example of α -pruning

- $5 < 7$, so update vals
- Can update min and beta



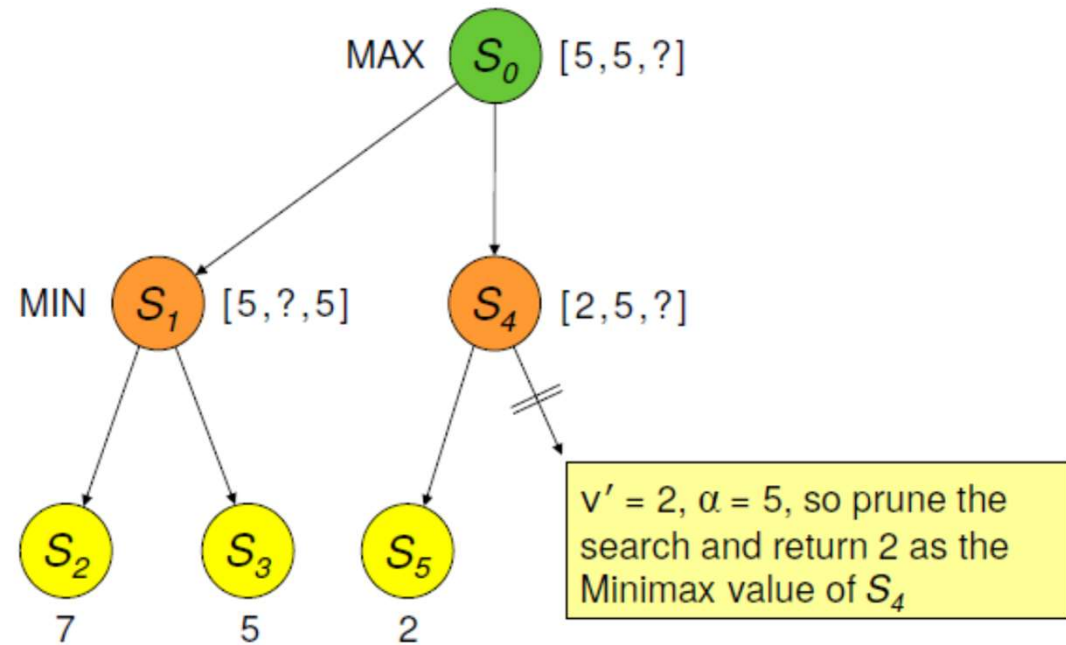
Example of α -pruning

- Similar, except start with max
 - i.e. max goes first



Example of α -pruning

- Similar, except start with max
 - i.e. max goes first



How Effective is Pruning?

- The effectiveness of pruning is highly dependent on the order in which successors are examined
 - Optimal if optimal branch examined first!
- In the best case, we can get $O(b^{m/2})$ rather than $O(b^m)$
 - Moves, rather than depth
- This means that alpha-beta can look ahead roughly twice as far as minimax in the same amount of time
- However, we can still only apply this to small games, chess and 5 x 5 dots and boxes are still out of reach!
- Next week: how to cope if you've only got a limited amount of time to make your decision...

Summary

- We've looked at two player zero-sum games of perfect information
- In these games, we can do a game tree search
- We use minimax and alpha beta pruning to search the game tree to find the optimal move