



CS310 – AI Foundations

Andrew Abel

March 2024

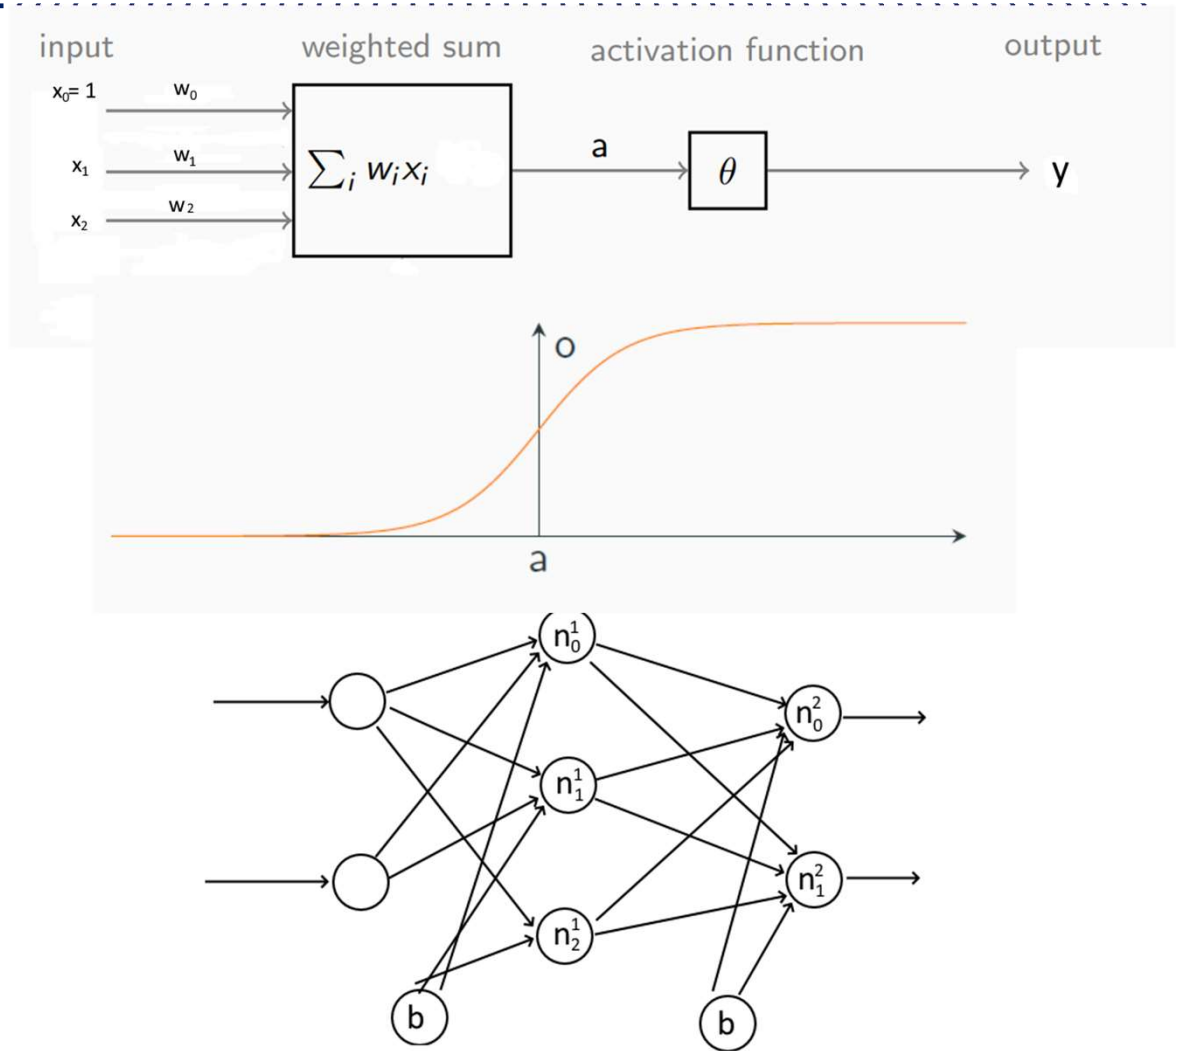

Week 8: Real Network Modelling

Welcome!

- Previous topic was very heavy!
 - A lot of gradient descent, backpropagation etc.
 - Now we will introduce the bigger picture
- Tensorflow
- Creating a model

Previously

- We had our perceptron
- Went from binary step to sigmoid activation
- Built a fully connected feedforward neural network
- Performed back propagation with gradient descent

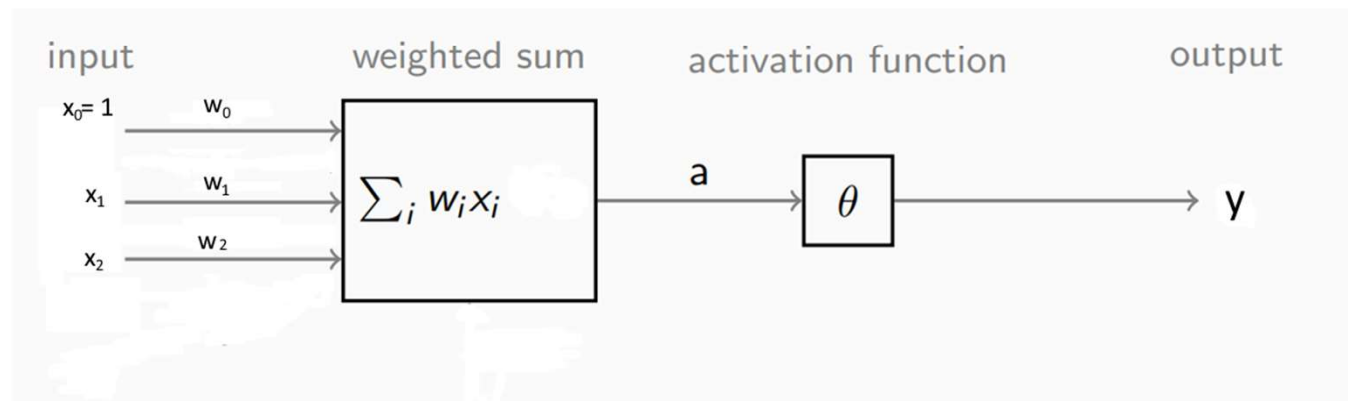
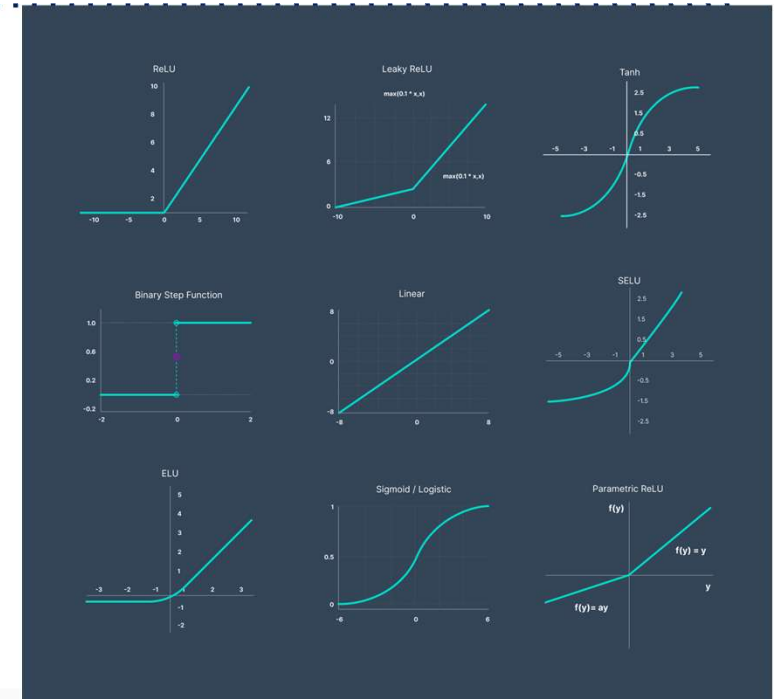


Is this the state-of-the-art?

- Much of that is because it is good for teaching
 - Sigmoid is easily differentiable manually
- Back propagation is the “classic” approach
 - Still in use today, but many variations!
- A lot of different activation functions
 - Sigmoid is no longer the recommended approach
- There are even different types of neurons


Activation Upgrades

- Still widely used!
- However, binary step and sigmoid are not so widely used
- Many different activation functions advised
- Discussed in previous lecture
- <https://www.v7labs.com/blog/neural-networks-activation-functions>



ReLU

- Rectified Linear Unit
- $f(x) = \max(0, x)$
- It looks linear, but does have a derivative
- Currently very fashionable
- Recommended for use in hidden layers by default
- Does not activate for negative values
 - More efficient
 - Can mean some weights and inputs are never trained (dead neurons)





Activation Function Recommendations

- Hidden layers recommendations
 - **Convolutional Neural Network (CNN)**: ReLU activation function.
 - **Recurrent Neural Network**: Tanh and/or Sigmoid activation function.
- Output layer recommendations
 - **Regression** - Linear Activation Function
 - **Binary Classification** - Sigmoid/Logistic Activation Function
 - **Multiclass Classification** - Softmax
 - **Multilabel Classification** - Sigmoid
- Trial and error works too!


Training and Optimisation

- Gradient Descent was the “standard”
- But many variations exist
 - Momentum, Adam etc.
- One of the mostly widely used is **Adam (Adaptive Moment Estimation)**
 - Generally a good choice!
- Not going to discuss here
- <https://towardsdatascience.com/optimizers-for-training-neural-network-59450d71caf6>



Fully Connected Networks

- Still extremely valuable and current
- Fully connected layers are known as dense layers
- These take an input, and feed it forward
- Most models will include at least 1 fully connected layer



Tensorflow for Machine Learning in Python

Building Your Own

- Python is a recommended language for machine learning
- There are two main libraries Tensorflow and Pytorch
 - Other libraries are available
- Tensorflow is enormous, we are not going to go through it all!
 - <https://www.tensorflow.org/>
- We will walk through creating a very simple network

Creating your Environment

- You will likely need to add several packages to your Python install
- tensorflow – the machine learning package
- pandas – Python Data Analysis package
- sklearn – Another machine learning package
- Keras – Python interface for machine learning
 - now effectively tied together with Tensorflow
- If running this on a laptop, you WILL get error/warning messages, this is designed for GPU processing, not CPU processing, but it will still work

```
1 # check version
2 import tensorflow
3 print(tensorflow.__version__)
```

Stages of creating a machine learning model

1. Prepare Dataset
2. Define the model.
 - Select layers, activations, etc.
3. Compile the model.
 - Creates the model
4. Fit the model.
 - Training process
5. Evaluate the model.
 - Test training performance
6. Make predictions.

Prepare Dataset

- Sign up for Kaggle
 - <https://www.kaggle.com/datasets?tags=13207-Computer+Vision>
 - One of the best resources for datasets
- Split into Training and Test sets (if not already done)
- You may want to normalise
- Read into Python
- Store as a pickle
 - Research this!

```
import numpy as np
data_path = "./"
train_data = np.loadtxt(data_path + "mnist_train.csv", delimiter=",")
test_data = np.loadtxt(data_path + "mnist_test.csv", delimiter=",")
```

Prepare Dataset II

- You may need to do some work to identify labels
- Also, if just one dataset, you may want to split into training and test sets

```
X_train, X_test, y_train, y_test
    = train_test_split(X, y, test_size=0.33)
```

 - Not needed for MNIST dataset
 - Can use `train_test_split` sklearn function
 - `Read_csv` function may also be useful
- Do similar work to last week, define training and test sets and labels
 - Do not need bias
 - Might want to define your lists as numpy arrays, then we can use “shape”

```
train_input = np.array([np.array(d[1:]) for d in train_data ])
```

Define the Model

- Select model you need, then choose architecture
 - define the layers of the model
 - configure each layer with a number of nodes and activation functions
 - connect the layers together into a cohesive model.
 - Tensorflow API useful
 - https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense
 - Also the guide
 - <https://machinelearningmastery.com/tensorflow-tutorial-deep-learning-with-tf-keras/>

```
# groups models together into a sequential model  
model = Sequential()
```


Define the Model

- Example model, create a sequential model (the simplest type)
- We create a fully connected hidden layer which receives input from our input (so 784 features)
 - Bias is added automatically
 - Has 4 neurons
 - Activation functions are sigmoid shaped
- A second layer is created, and in this case we use 10 neurons (i.e. one for each number we are trying to predict)
 - We use sigmoid again, but 'softmax' might be better
- Can create as many layers as we want
 - The guide helps with options!
 - https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense

```
# Step 2, define model
model = Sequential()
model.add(Dense(4, activation='sigmoid', input_shape=(n_features,)))
model.add(Dense(10, activation='sigmoid'))
```

Compile the Model

- Compiling the model requires that you first select a loss function that you want to optimize, such as mean squared error or cross-entropy.
- Requires that you select an algorithm to perform the optimization procedure
 - Can use stochastic gradient descent
 - more modern variations, such as Adam are available
- The optimizer can be specified as a string for a known optimizer class, e.g. *'sgd'* for stochastic gradient descent, or you can configure an instance of an optimizer class and use that.
 - Note here, can define learning rate and momentum (additional property)

```
# compile the model
```

```
opt = SGD(learning_rate=0.01, momentum=0.9)
```

```
model.compile(optimizer=opt, loss='binary_crossentropy')
```

Compile the Model

- The three most common loss functions:
 - `'binary_crossentropy'` for binary classification.
 - `'sparse_categorical_crossentropy'` for multi-class classification.
 - `'mse'` (mean squared error) for regression.
- More here
 - https://www.tensorflow.org/api_docs/python/tf/keras/optimizers
- Here we use `sgd`
- Test the loss with `crossentropy` (I had problems with `MSE` for this problem!)
- Use `accuracy` as our measure

```
# Step 3, compile model
model.compile(optimizer='sgd', loss='SparseCategoricalCrossentropy',
metrics=['accuracy'])
model.summary()
```

Compile the model

- A number of approaches you can use
 - Experiment and see!
- <https://analyticsindiamag.com/guide-to-tensorflow-keras-optimizers/>
- <https://python.plainenglish.io/mastering-optimizers-with-tensorflow-a-deep-dive-into-efficient-model-training-81c58c630ef1>

```
model.compile(optimizer='adam', loss='categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.compile(optimizer='sgd', loss='SparseCategoricalCrossentropy',  
              metrics=['accuracy'])
```

```
opt = SGD(learning_rate=0.01, momentum=0.9)model.compile(optimizer=opt,  
                  loss='binary_crossentropy')  
model.compile(optimizer=opt, loss='binary_crossentropy')
```

Fit the Model

- 'Fitting' refers to training
 - 'Fit' the model to the data

```
model.fit(train_input, train_label, epochs=5, batch_size=32,  
verbose=1)
```

- Need to select the training config
 - Number of epochs (number of times to train)
 - Batch size (samples used to estimate the error)
- Training tries to minimise the loss function, and backpropagates the model
- Slowest part of machine learning. Can take hours, days, even weeks!
- By setting "verbose" to 1, can see progress
 - Set to 0 to get no display

Fit the Model

- Slowest part of machine learning. Can take hours, days, even weeks!
- More epochs means it might fit the model better, but might cause **overfitting**
- By setting “verbose” to 1, can see progress
 - Set to 0 to get no display

```
Epoch 1/5
1875/1875 [=====] - 1s 709us/step - loss: 1.9657 - accuracy: 0.2945
Epoch 2/5
1875/1875 [=====] - 1s 613us/step - loss: 1.7178 - accuracy: 0.3834
Epoch 3/5
1875/1875 [=====] - 1s 598us/step - loss: 1.6049 - accuracy: 0.3994
Epoch 4/5
1875/1875 [=====] - 1s 613us/step - loss: 1.5169 - accuracy: 0.4136
Epoch 5/5
1875/1875 [=====] - 1s 598us/step - loss: 1.4732 - accuracy: 0.4106
```

Evaluate the Model

- You should use separate data for testing your model
 - we should get an unbiased estimate of the performance of the model when making predictions on new data.
- Much faster than training, as you do not need to adjust the weights
- From an API perspective, this involves calling a function with the holdout dataset and getting a loss and perhaps other metrics that can be reported.

```
# Step 5, evaluate the model
loss, acc = model.evaluate(test_input, test_label, verbose=0)
print('Test Accuracy: %.3f' % acc)
```

Predictions!

- The fun part of the model
 - How well can it classify your input?
- For this one, we are looking to see how well it classifies, so which one of the 10 cases
- Other problems may want to see a different output

```
print("actual: ", test_label[22])
row = np.array(test_input[22])[None, ...]
yhat = model.predict(row)
print('Predicted: %s (class=%d)' % (yhat,
np.argmax(yhat)))
fig = plt.figure()
plt.plot(yhat[0])
plt.show()
```


Other tips

- When working with this code, it can be easy to mess up configurations
- Create a separate virtual environment (venv)
 - This gives you a clean environment for you to install packages
 - Doesn't install them elsewhere
- If you do not have a GPU, then you may get warning messages about using tensorflow
 - It will still work, but a little slower
- A lot of the work will be trial and error. Experiment with different configurations

And that is your first model!

- You should improve on this for your week 9 lab
- Experiment with settings
 - Reading input data
 - Creating dense (fully connected layers)
 - Generating output
 - Testing